

---

## **ARN - Report - Labo04**

Anthony Coke, Guilain Mbayo, Mehdi Salhi



May 11, 2022

## Contents

<b>Learning algorithm</b>	<b>3</b>
<b>Model Complexity</b>	<b>3</b>
<b>Deep Neural Networks</b>	<b>4</b>
<b>Tests</b>	<b>4</b>

## Learning algorithm

1. What is the learning algorithm being used to optimize the weights of the neural networks? What are the parameters (arguments) being used by that algorithm? What cost function is being used ? please, give the equation(s)

MLP\_from\_raw\_data.ipynb

The algorithm used is RMSprop.

The arguments used by this algorithm are: - Learning rate A Tensor, floating point value, or a schedule that is a `tf.keras.optimizers.schedules.LearningRateSchedule`, or a callable that takes no arguments and returns the actual value to use. The learning rate. Defaults to 0.001. - rho: Discounting factor for the history/coming gradient. Defaults to 0.9. - momentum: A scalar or a scalar Tensor. Defaults to 0.0. - epsilon: A small constant for numerical stability. This epsilon is “epsilon hat” in the Kingma and Ba paper (in the formula just before Section 2.1), not the epsilon in Algorithm 1 of the paper. Defaults to 1e-7. - centered: Boolean. If True, gradients are normalized by the estimated variance of the gradient; if False, by the uncentered second moment. Setting this to True may help with training, but is slightly more expensive in terms of computation and memory. Defaults to False. - name: Optional name prefix for the operations created when applying gradients. Defaults to “RMSprop”. - \*\*kwargs: keyword arguments. Allowed arguments are clipvalue, clipnorm, global\_clipnorm. If clipvalue (float) is set, the gradient of each weight is clipped to be no higher than this value. If clipnorm (float) is set, the gradient of each weight is individually clipped so that its norm is no higher than this value. If global\_clipnorm (float) is set the gradient of all weights is clipped so that their global norm is no higher than this value.

The used cost function is the categorical crossentropy function. It's equation is:

$$\text{Loss} = - \sum_{i=1}^{\text{output size}} y_i \cdot \log \hat{y}_i$$

**Figure 1:** ARN-Labo04-CrossEntrEquation

## Model Complexity

2. Model complexity: for each experiment (shallow network learning from raw data, shallow network learning from features, CNN, and Fashion MNIST), select a neural network topology and describe the inputs, indicate how many are they, and how many outputs. Compute

the number of weights of each model (e.g., how many weights between the input and the hidden layer, how many weights between each pair of layers, biases, etc..) and explain how do you get to the total number of weights.

MLP\_from\_raw\_data.ipynb Inputs: Outputs: Activation function: tanh Activation function for output layer: softmax Neurons in hidden layer: 250 Batch size: 4096 Dropout: 0.5 Number of epoch: 150

MLP\_from\_HOG.ipynb

CNN.ipynb

Fashion\_MNIST.ipynb

## Deep Neural Networks

3. Do the deep neural networks have much more “capacity” (i.e., do they have more weights?) than the shallow ones? explain with one example

The deep neural network have more hidden layer than the shallow ones, but it doesn't necessary mean that it has more neurons in it. For exemple, in this lab we use 300 neurons in the hidden layer for the shallows network (raw\_data and HOG), against only 25 neurons for the deep one (CNN). The deep neural networks have more capacity, because they usually need less components to achieve the same goal or better than a shallow neural network. If we compare the weights of each model, the shallow one will have more weight than the deep one. For exemple, a model with 2 entries, 6 neurons in one hidden layer and 2 output, we get  $2 * 6 + 6 * 2 = 24$  links that have each their weight. For the same model but with 3 hidden layers, we got  $2 * 2 + 2 * 2 + 2 * 2 + 2 * 2 = 16$  links, and so 16 weights.

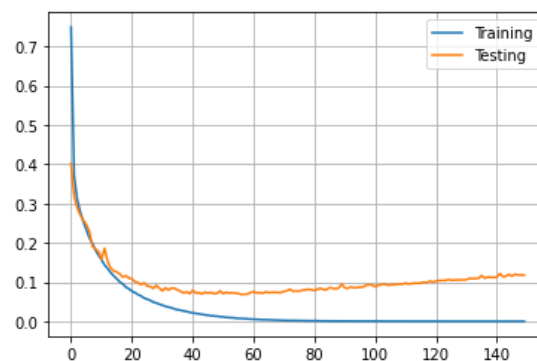
## Tests

4. Test every notebook for at least three different meaningful cases (e.g., for the MLP exploiting raw data, test different models varying the number of hidden neurons, for the feature-based model, test pix\_p\_cell 4 and 7, and number of orientations or number of hidden neurons, for the CNN, try different number of neurons in the feed-forward part) describe the model and present the performance of the system (e.g., plot of the evolution of the error, final evaluation scores and confusion matrices). Comment the differences in results. Are there


particular digits that are frequently confused?

MLP\_from\_raw\_data.ipynb

Test score: 0.11751019209623337  
Test accuracy: 0.9818999767303467




**Figure 2:** ARN-RAW-Plot-tanh-softmax\_Batch2048\_NoDropout\_Epoch150




**Figure 3:** ARN-RAW-ConfMat-tanh-softmax\_Batch2048\_NoDropout\_Epoch150

We can see in this experiment that there's clearly an overfitting.

Test score: 0.0734260305762291  
Test accuracy: 0.9796000123023987




**Figure 4:** ARN-RAW-Plot-tanh-softmax\_Batch2048\_Dropout\_Epoch150

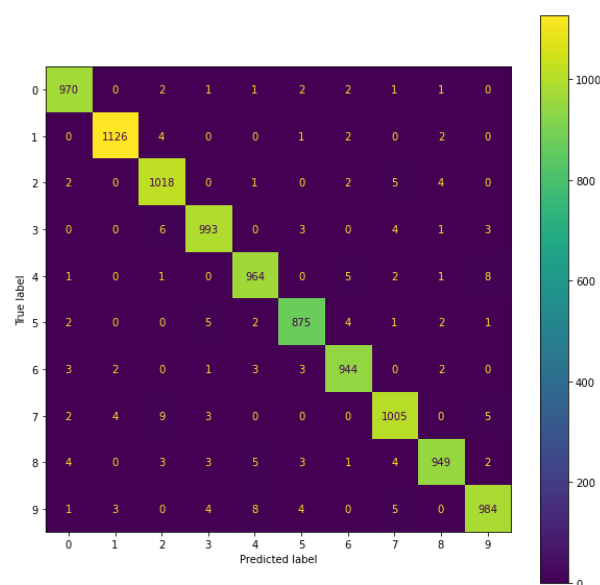


**Figure 5:** ARN-RAW-ConfMat-tanh-softmax\_Batch2048\_Dropout\_Epoch150

Test score: 0.06564721465110779  
Test accuracy: 0.9828000068664551

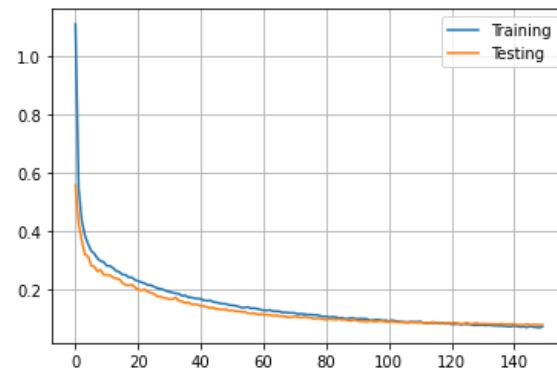


**Figure 6:** ARN-RAW-Plot-sigmoid-softmax\_Batch2048\_Dropout\_Epoch150

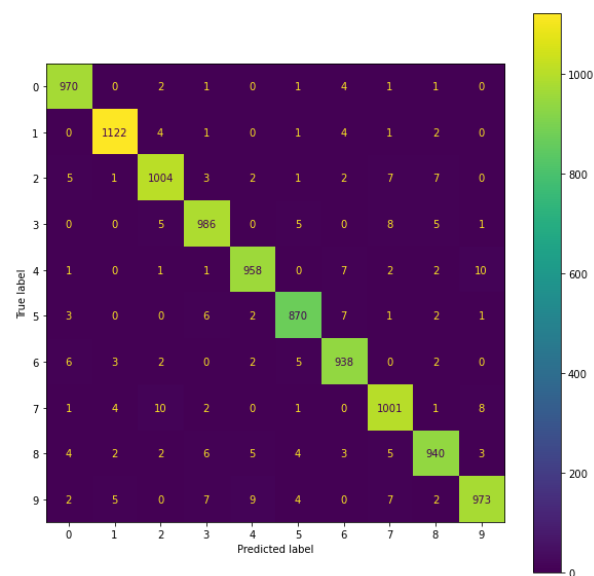


**Figure 7:** ARN-RAW-ConfMat-sigmoid-softmax\_Batch2048\_Droptout\_Epoch150

Test score: 0.08130748569965363  
Test accuracy: 0.9761999845504761



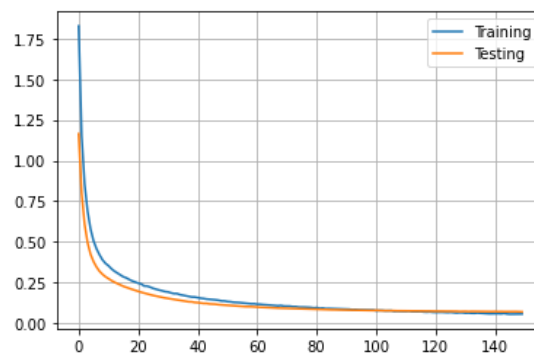
**Figure 8:** ARN-RAW-Plot-tanh-softmax-Neur250\_Batch4096\_Dropout\_Epoch150




**Figure 9:** ARN-RAW-ConfMat-tanh-softmax-Neur250\_Batch4096\_Dropout\_Epoch150



Test score: 0.06945059448480606  
Test accuracy: 0.9797000288963318

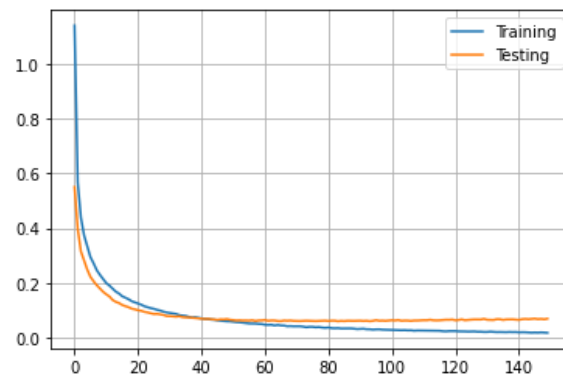


**Figure 10:** ARN-RAW-Plot-sigmoid-softmax-Neur250\_Batch4096\_Dropout\_Epoch150

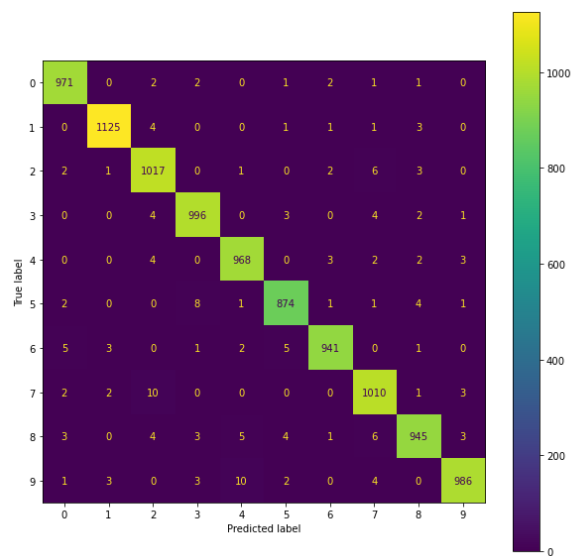


**Figure 11:** ARN-RAW-ConfMax-sigmoid-softmax-Neur250\_Batch4096\_Dropout\_Epoch150

Test score: 0.06911627948284149  
Test accuracy: 0.983299970626831

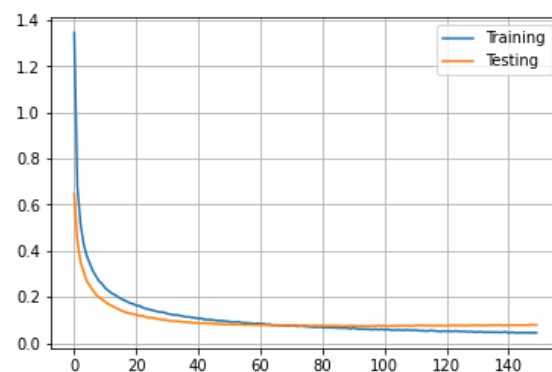


**Figure 12:** ARN-RAW-Plot-relu\_softmax-Neur250\_Batch4096\_Dropout\_Epoch150




**Figure 13:** ARN-RAW-ConfMat-relu\_softmax-Neur250\_Batch4096\_Dropout\_Epoch150

Test score: 0.07871479541063309  
Test accuracy: 0.9796000123023987

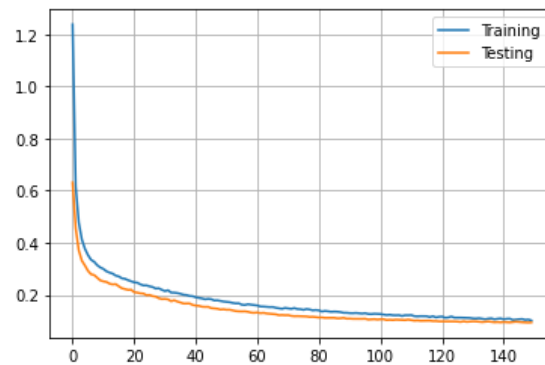


**Figure 14:** ARN-RAW-Plot-relu-softmax-Neur150\_Batch4096\_Dropout\_Epoch150

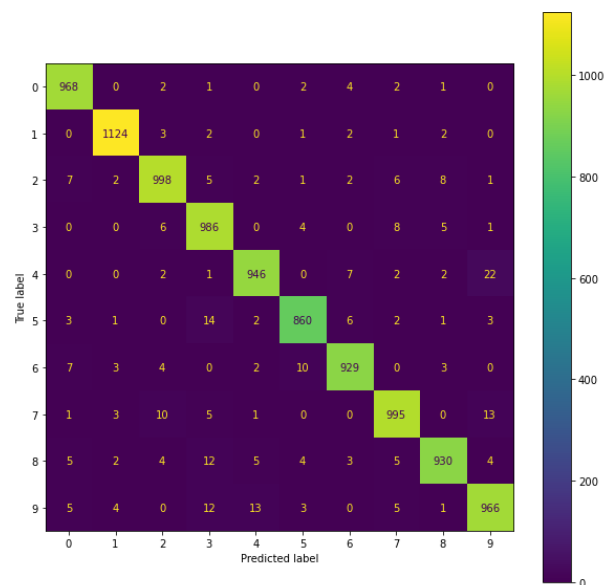


**Figure 15:** ARN-RAW-ConfMat-relu-softmax-Neur150\_Batch4096\_Dropout\_Epoch150

Test score: 0.09491664171218872  
Test accuracy: 0.9702000021934509

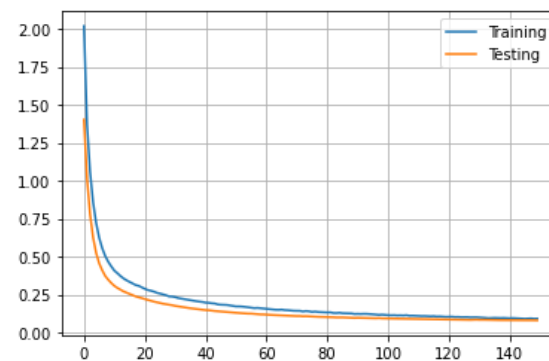


**Figure 16:** ARN-RAW-Plot-tanh-softmax-Neur150\_Batch4096\_Dropout\_Epoch150




**Figure 17:** ARN-RAW-ConfMat-tanh-softmax-Neur150\_Batch4096\_Dropout\_Epoch150

Test score: 0.08042246848344803  
Test accuracy: 0.9757999777793884



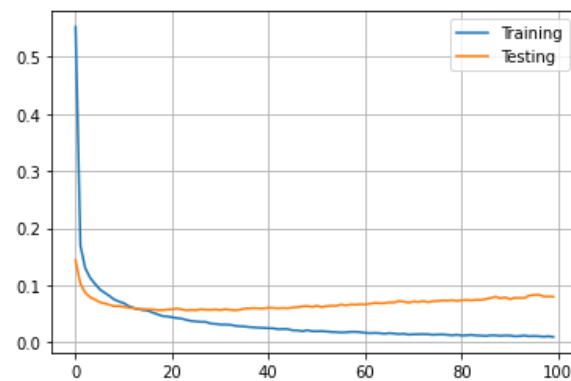
**Figure 18:** ARN-RAW-Plot-sigmoid-softmax-Neur150\_Batch4096\_Dropout\_Epoch150



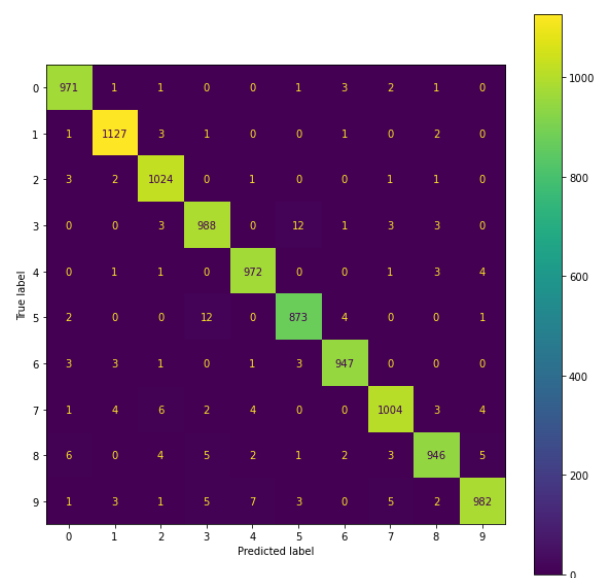
**Figure 19:** ARN-RAW-ConfMat-sigmoid-softmax-Neur150\_Batch4096\_Dropout\_Epoch150

MLP\_from\_HOG.ipynb

Test score: 0.08003607392311096  
Test accuracy: 0.9833999872207642




**Figure 20:** ARN-HOG-Plot-relu-softmax-Neur200\_Batch512\_Dropout\_Epoch100




**Figure 21:** ARN-HOG-ConfMat-relu-softmax-Neur200\_Batch512\_Dropout\_Epoch100

Test score: 0.08840184658765793  
Test accuracy: 0.97560004196167



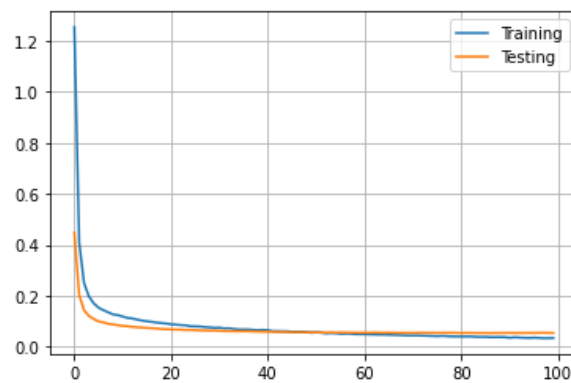
**Figure 22:** ARN-HOG-Plot-tanh-softmax-Neur200-Batch512\_Dropout\_Epoch100



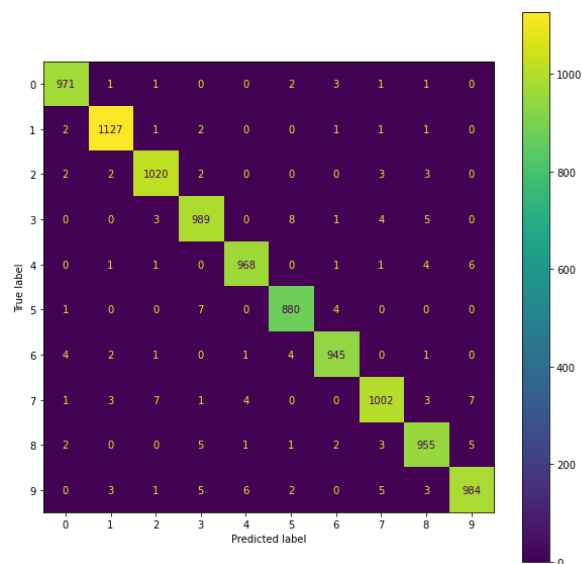
**Figure 23:** ARN-HOG-ConfMat-tanh-softmax-Neur200\_Batch512\_Dropout\_Epoch100



Test score: 0.05404368415474892  
Test accuracy: 0.9840999841690063

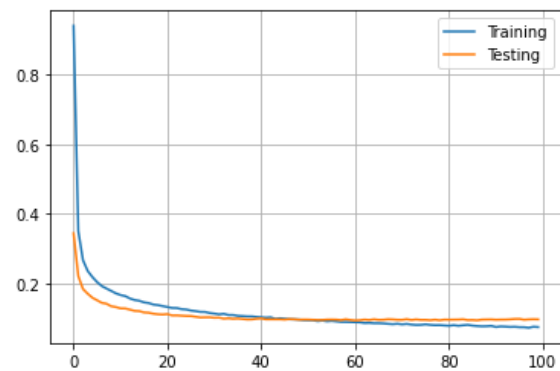


**Figure 24:** ARN-HOG-Plot-sigmoid-softmax-Neur200\_Batch512\_Dropout\_Epoch100

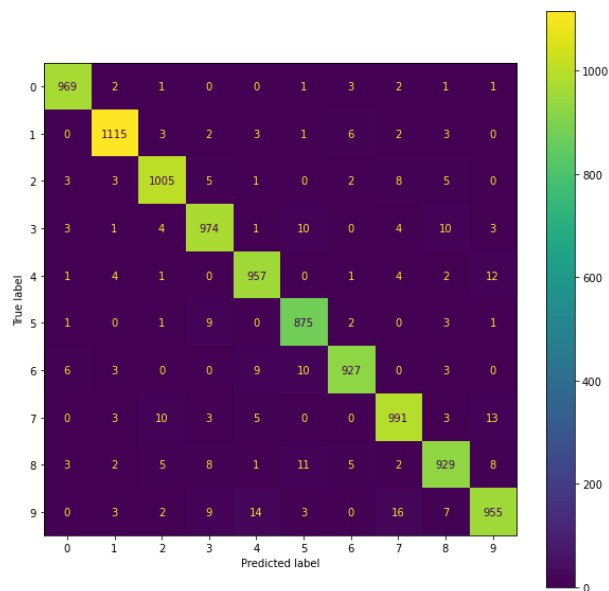


**Figure 25:** ARN-HOG-ConfMat-sigmoid-softmax-Neur200\_Batch512\_Dropout\_Epoch100

Test score: 0.09604205936193466  
Test accuracy: 0.9696999788284302




**Figure 26:** ARN-HOG-Plot-relu-Pixel7\_Neur200\_Batch512\_Dropout\_Epoch100

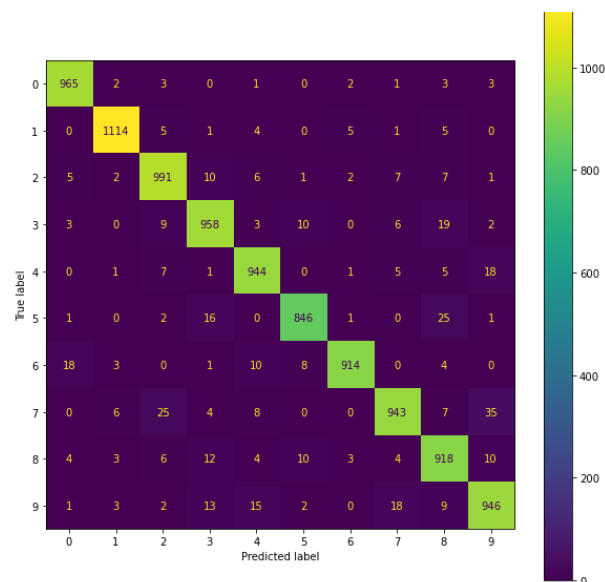


**Figure 27:** ARN-HOG-ConfMat-relu-Pixel7\_Neur200\_Batch512\_Dropout\_Epoch100

Test score: 0.13775815069675446  
Test accuracy: 0.9538999795913696




**Figure 28:** ARN-HOG-Plot-tanh-Pixel7\_Neur200\_Batch512\_Dropout\_Epoch100




**Figure 29:** ARN-HOG-ConfMat-tanh-Pixel7\_Neur200\_Batch512\_Dropout\_Epoch100

Test score: 0.09379129111766815  
Test accuracy: 0.9699000120162964

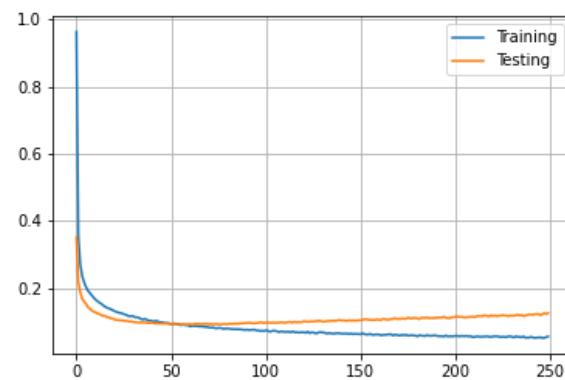


**Figure 30:** ARN-HOG-Plot-sigmoid-Pixel7\_Neur200\_Batch512\_Dropout\_Epoch250



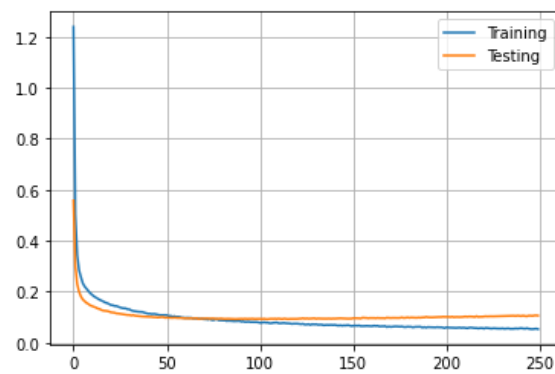
**Figure 31:** ARN-HOG-ConfMat-sigmoid-Pixel7\_Neur200\_Batch512\_Dropout\_Epoch250

Test score: 0.12630851566791534  
Test accuracy: 0.9679999947547913



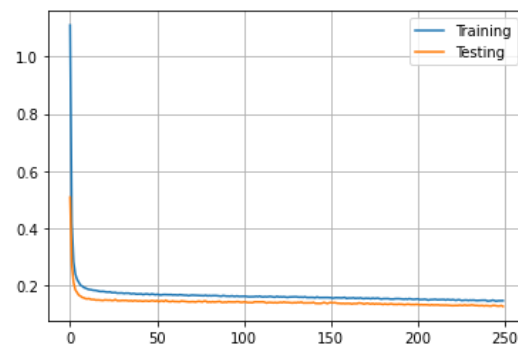
**Figure 32:** ARN-HOG-Plot-relu-Pixel7\_Neur200\_Batch512\_Dropout\_Epoch250

Test score: 0.10599838942289352  
Test accuracy: 0.9700999855995178



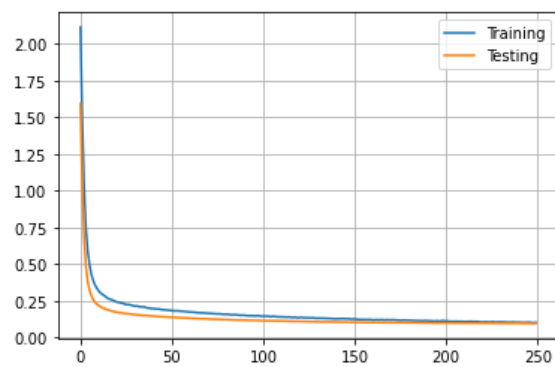
**Figure 33:** ARN-HOG-Plot-relu-Pixel7\_Neur200\_Batch1024\_Dropout\_Epoch250

Test score: 0.1260688304901123  
Test accuracy: 0.957599975204468




**Figure 34:** ARN-HOG-Plot-tanh-Pixel7\_Neur200\_Batch1024\_Dropout\_Epoch250

Test score: 0.0942075327038765  
Test accuracy: 0.9684000015258789

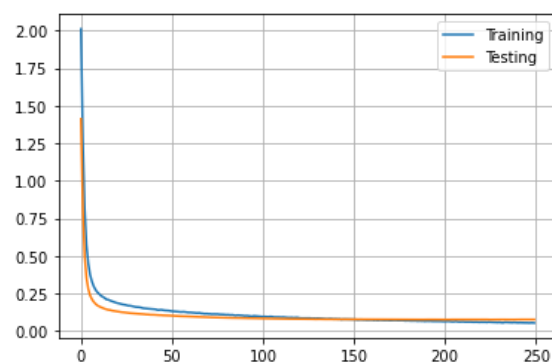


**Figure 35:** ARN-HOG-Plot-sigmoid-Pixel7\_Neur200\_Batch1024\_Dropout\_Epoch250




**Figure 36:** ARN-HOG-ConfMat-sigmoid-Pixel7\_Neur200\_Batch1024\_Dropout\_Epoch250


Test score: 0.0766327977180481  
 Test accuracy: 0.9750999808311462



**Figure 37:** ARN-HOG-Plot-sigmoid-Pixel7\_Or16\_Neur200\_Batch1024\_Dropout\_Epoch250




**Figure 38:** ARN-HOG-Plot-relu-Pixel7\_Or16\_Neur200\_Batch1024\_Dropout\_Epoch250




**Figure 39:** ARN-HOG-Plot-tanh-Pixel7\_Or16\_Neur200\_Batch1024\_Dropout\_Epoch200






Test score: 0.09441999346017838  
Test accuracy: 0.9700999855995178




**Figure 40:** ARN-HOG-Plot-relu-Pixel7\_Or8\_Neur150\_Batch1024\_Dropout\_Epoch200




**Figure 41:** ARN-HOG-ConfMat-relu-Pixel7\_Or8\_Neur150\_Batch1024\_Dropout\_Epoch200

Test score: 0.1024104580283165  
Test accuracy: 0.9660999774932861




**Figure 42:** ARN-HOG-Plot-sigmoid-Pixel7\_Or8\_Neur150\_Batch1024\_Dropout\_Epoch200




**Figure 43:** ARN-HOG-ConfMat-sigmoid-Pixel7\_Or8\_Neur150\_Batch1024\_Dropout\_Epoch200

Test score: 0.1345009058713913  
Test accuracy: 0.9556000232696533




**Figure 44:** ARN-HOG-Plot-tanh-Pixel7\_Or8\_Neur150-Batch1024\_Dropout\_Epoch200

Test score: 0.10055689513683319  
Test accuracy: 0.9699000120162964




**Figure 45:** ARN-HOG-Plot-relu-Pixel7\_Or8\_Neur250\_Batch1024\_Dropout\_Epoch200

Test score: 0.09292542189359665  
Test accuracy: 0.9690999984741211



**Figure 46:** ARN-HOG-Plot-sigmoid-Pixel7\_Or8\_Neur250\_Batch1024\_Dropout\_Epoch200

Test score: 0.13129283487796783  
Test accuracy: 0.9555000066757202



**Figure 47:** ARN-HOG-Plot-tanh-Pixel7\_Or8\_Neur250\_Batch1024\_Dropout\_Epoch200

CNN.ipynb