
ARN - Report - Labo04

Anthony Coke, Guilain Mbayo, Mehdi Salhi



May 12, 2022

Contents

Learning algorithm	3
Model Complexity	3
Deep Neural Networks	5
Tests	6
Conclusion - RAW	15
Conclusion - HOG	27
Conclusion - CNN	34
Conclusion - CNN Fashion MNIST	34

Learning algorithm

1. What is the learning algorithm being used to optimize the weights of the neural networks? What are the parameters (arguments) being used by that algorithm? What cost function is being used ? please, give the equation(s)

MLP_from_raw_data.ipynb

The algorithm used is RMSprop.

The arguments used by this algorithm are: - Learning rate A Tensor, floating point value, or a schedule that is a `tf.keras.optimizers.schedules.LearningRateSchedule`, or a callable that takes no arguments and returns the actual value to use. The learning rate. Defaults to 0.001. - rho: Discounting factor for the history/coming gradient. Defaults to 0.9. - momentum: A scalar or a scalar Tensor. Defaults to 0.0. - epsilon: A small constant for numerical stability. This epsilon is “epsilon hat” in the Kingma and Ba paper (in the formula just before Section 2.1), not the epsilon in Algorithm 1 of the paper. Defaults to 1e-7. - centered: Boolean. If True, gradients are normalized by the estimated variance of the gradient; if False, by the uncentered second moment. Setting this to True may help with training, but is slightly more expensive in terms of computation and memory. Defaults to False. - name: Optional name prefix for the operations created when applying gradients. Defaults to “RMSprop”. - **kwargs: keyword arguments. Allowed arguments are clipvalue, clipnorm, global_clipnorm. If clipvalue (float) is set, the gradient of each weight is clipped to be no higher than this value. If clipnorm (float) is set, the gradient of each weight is individually clipped so that its norm is no higher than this value. If global_clipnorm (float) is set the gradient of all weights is clipped so that their global norm is no higher than this value.

The used cost function is the categorical crossentropy function. It's equation is:

$$\text{Loss} = - \sum_{i=1}^{\text{outputsize}} y_i \cdot \log(\hat{y}_i)$$

Model Complexity

2. Model complexity: for each experiment (shallow network learning from raw data, shallow network learning from features, CNN, and Fashion MNIST), select a neural network topology and describe the inputs, indicate how many are they, and how many outputs. Compute the number of weights of each model (e.g., how many weights between the input and the hidden layer, how many weights between each pair of layers, biases, etc..) and explain how do you get to the total number of weights.

MLP_from_raw_data.ipynb

Inputs: 784, which are each pixels in a picture

Outputs: 10 classes (numbers between 0 and 9)

Activation function: tanh

Activation function for output layer: softmax

Neurons in hidden layer: 250

Batch size: 4096

Dropout: 0.5

Number of epoch: 150

The model has 784 inputs, 1 hidden layer that contains 250 neurons and 10 outputs. The number of weights between the inputs and the hidden layer is $784 * 250 = 196000$. The number of weights between the hidden layer and the outputs is $250 * 10 = 2500$. The total number of weights is 198500.

MLP_from_HOG.ipynb

Inputs: 392

Outputs: 10 classes (numbers between 0 and 9)

Activation function: sigmoid

Activation function for output layer: softmax

Neurons in hidden layer: 200

Batch size: 1024

pixel per cell: 7

n_orientation: 16

number of epoch: 250 (but we could see that 150 is enough)

Dropout: 0.5

The model has 392 inputs, 1 hidden layer that contains 200 neurons and 10 outputs. The number of weights between the inputs and the hidden layer is $392 * 200 = 78400$. The number of weights between the hidden layer and the outputs is $200 * 10 = 2000$. The total number of weights is 80400.

CNN.ipynb

Inputs: 144 for the MLP and a 28 X 28 pixels image as input for the features extraction part. The 28 X 28 image is passed through 3 features extraction layers. The first one take work with 28 X 28 images, the second with 14 X 14 images and the third with 7 X 7 images. The result is then flattened into a vector of 144 values and passed to the classification part.

Outputs: 10 classes (numbers between 0 and 9)

Activation function: sigmoid

Activation function for output layer: softmax

Neurons in hidden layer: 25

Batch size: 256

Number of epoch: 50

The model has 144 inputs, 1 hidden layer that contains 25 neurons and 10 outputs. The number of weights between the inputs and the hidden layer is $144 * 25 = 3600$. The number of weights between the hidden layer and the outputs is $25 * 10 = 2500$. The total number of weights is 6100.

Fashion_MNIST.ipynb

Inputs:

Outputs: 10 classes (numbers between 0 and 9)

Activation function:

Activation function for output layer: softmax

Neurons in hidden layer:

Batch size:

Number of epoch:

Deep Neural Networks

3. Do the deep neural networks have much more “capacity” (i.e., do they have more weights?) than the shallow ones? explain with one example

The deep neural network have more hidden layer than the shallow ones, but it doesn't necessary mean that it has more neurons in it. For exemple, in this lab we use 300 neurons in the hidden layer for the shallows network (raw_data and HOG), against only 25 neurons for the deep one (CNN). The deep neural networks have more capacity, because they usually need less components to achieve the same

goal or better than a shallow neural network. If we compare the weights of each model, the shallow one will have more weight than the deep one. For example, a model with 2 entries, 6 neurons in one hidden layer and 2 output, we get $2 * 6 + 6 * 2 = 24$ links that have each their weight. For the same model but with 3 hidden layers, we got $2 * 2 + 2 * 2 + 2 * 2 + 2 * 2 = 16$ links, and so 16 weights.

Tests

4. Test every notebook for at least three different meaningful cases (e.g., for the MLP exploiting raw data, test different models varying the number of hidden neurons, for the feature-based model, test pix_p_cell 4 and 7, and number of orientations or number of hidden neurons, for the CNN, try different number of neurons in the feed-forward part) describe the model and present the performance of the system (e.g., plot of the evolution of the error, final evaluation scores and confusion matrices). Comment the differences in results. Are there particular digits that are frequently confused?

MLP_from_raw_data.ipynb

Model:

- Activation function: tanh
- Neurons: 300
- Dropout: -
- Batch size: 2048
- Epochs: 150

Test score: 0.11751019209623337
Test accuracy: 0.9818999767303467




Figure 1: ARN-RAW-Plot-tanh-softmax_Batch2048_NoDropout_Epoch150




Figure 2: ARN-RAW-ConfMat-tanh-softmax_Batch2048_NoDropout_Epoch150

We first tried to train a model without a dropout. As the graph shows there is a problem with overfitting. Surprisingly the confusion matrix is not so bad, the wrong classifications are low, but there are a lot.

Model:

- Activation function: tanh
- Neurons: 300
- Dropout: 0.5
- Batch size: 2048
- Epochs: 150

Test score: 0.0734260305762291
 Test accuracy: 0.9796000123023987




Figure 3: ARN-RAW-Plot-tanh-softmax_Batch2048_Dropout_Epoch150




Figure 4: ARN-RAW-ConfMat-tanh-softmax_Batch2048_Dropout_Epoch150

We added a dropout and the result already improved. The gap has decreased but at the beginning the test set looks too easy. The confusion matrix shows higher errors but no number really stands out.

Model:

- Activation function: tanh
- Neurons: 250
- Dropout: 0.5
- Batch size: 4096
- Epochs: 150

Test score: 0.08130748569965363
Test accuracy: 0.9761999845504761




Figure 5: ARN-RAW-Plot-tanh-softmax-Neur250_Batch4096_Dropout_Epoch150




Figure 6: ARN-RAW-ConfMat-tanh-softmax-Neur250_Batch4096_Dropout_Epoch150

With fewer neurons and a batch size of 4096, we get a nice result. The two curves are really close to each other. We begin to see numbers with two digits in the confusion matrix. Numbers 4 and 7 are the less good classified.

Model:

- Activation function: sigmoid
- Neurons: 250
- Dropout: 0.5
- Batch size: 4096
- Epochs: 150

Test score: 0.06945059448480606
 Test accuracy: 0.9797000288963318




Figure 7: ARN-RAW-Plot-sigmoid-softmax-Neur250_Batch4096_Dropout_Epoch150




Figure 8: ARN-RAW-ConfMax-sigmoid-softmax-Neur250_Batch4096_Dropout_Epoch150

When changing the activation function to sigmoid the errors drop faster, the meeting point is similar to the last graph. The same numbers (4 and 7) are less good classified with two digits boxes.

Model:

- Activation function: tanh
- Neurons: 150
- Dropout: 0.5
- Batch size: 4096
- Epochs: 150

Test score: 0.09491664171218872
Test accuracy: 0.9702000021934509




Figure 9: ARN-RAW-Plot-tanh-softmax-Neur150_Batch4096_Dropout_Epoch150




Figure 10: ARN-RAW-ConfMat-tanh-softmax-Neur150_Batch4096_Dropout_Epoch150

Setting the number of neurons to 150 changes the moment when the two curves meet. With tanh we see that the testing set is too simple, the testing curve is always below the training one. More epochs are needed to find a stable ground. Too few neurons is not a good idea, it increases the amount of errors. We can see numbers up to 22. This model seems to have more difficulty to classify numbers from 4 to 9.

Model:

- Activation function: sigmoid
- Neurons: 150
- Dropout: 0.5
- Batch size: 4096
- Epochs: 150

Test score: 0.08042246848344803
 Test accuracy: 0.9757999777793884




Figure 11: ARN-RAW-Plot-sigmoid-softmax-Neur150_Batch4096_Dropout_Epoch150




Figure 12: ARN-RAW-ConfMat-sigmoid-softmax-Neur150_Batch4096_Dropout_Epoch150

The result is similar with the sigmoid function, more epochs are needed to find the meeting point. The errors are however less significant than with tanh in the confusion matrix. 4, 5, 7 and 10 are the only numbers with two digits errors.

Conclusion - RAW

Removing neurons seems to delay the moment when the two curves meet. We didn't find a model that stands out. Some models have a good graph, but some numbers have a 1-2% error in the confusion matrix. When computing the f1_score we obtained most of the time a score around 0.95 which is good. But this score doesn't really represent the model's behavior. Sometimes the graph clearly shows an overfitting but the f1_score is 0.97. So when choosing a model we should take everything into account, the graph, the f1-score, the confusion matrix.

MLP_from_HOG.ipynb

Model:

- Activation function: relu
- Neurons: 200
- Batch size: 512
- Dropout: 0.5
- Epochs: 100
- Pixels: 4
- Orientations: 8

Test score: 0.08003607392311096
Test accuracy: 0.9833999872207642




Figure 13: ARN-HOG-Plot-relu-softmax-Neur200_Batch512_Dropout_Epoch100




Figure 14: ARN-HOG-ConfMat-relu-softmax-Neur200_Batch512_Dropout_Epoch100

We tried to use the preconfigured model but with a batch size of 512 so it could go faster. The scores are not so bad, but the graph shows some overfitting starting from 10 epochs. The confusion matrix shows that number 3 is often assimilated as a 5 and number 5 as a 3. Those are the only values that strike out.

Model:

- Activation function: sigmoid
- Neurons: 200
- Batch size: 512
- Dropout: 0.5
- Epochs: 100
- Pixels: 4
- Orientations: 8

Test score: 0.05404368415474892
 Test accuracy: 0.9840999841690063




Figure 15: ARN-HOG-Plot-sigmoid-softmax-Neur200_Batch512_Dropout_Epoch100




Figure 16: ARN-HOG-ConfMat-sigmoid-softmax-Neur200_Batch512_Dropout_Epoch100

We decided to try different activation functions with the preconfigured model. This one shows the sigmoid function. As we can see this one behaves better than the previous and there's almost no overfitting. However we can see that the testing error is below the testing error until around 50 epochs

which might signify that the test set is a little bit too easy. The confusion matrix doesn't indicate any abnormalities.

Model:

- Activation function: tanh
- Neurons: 200
- Batch size: 512
- Dropout: 0.5
- Epochs: 100
- Pixels: 7
- Orientations: 8

Test score: 0.13775815069675446
Test accuracy: 0.9538999795913696




Figure 17: ARN-HOG-Plot-tanh-Pixel7_Neur200_Batch512_Dropout_Epoch100




Figure 18: ARN-HOG-ConfMat-tanh-Pixel7_Neur200_Batch512_Dropout_Epoch100

Those previous attempts were done with 4 pixels, we also tried 7 pixels. This example with the tanh function clearly depicts a test set that is too easy. The testing curves never overlaps the training one. We notice that the test score is really bad compared to the previous models. Of course the confusion matrix shows a similar result. This model doesn't seem to recognize the numbers between 2 and 9. In some cases we can see up to 35 wrong classification in one box.

Model:

- Activation function: sigmoid
- Neurons: 200
- Batch size: 512
- Dropout: 0.5
- Epochs: 250
- Pixels: 7
- Orientations: 8

Test score: 0.09379129111766815
 Test accuracy: 0.9699000120162964




Figure 19: ARN-HOG-Plot-sigmoid-Pixel7_Neur200_Batch512_Dropout_Epoch250




Figure 20: ARN-HOG-ConfMat-sigmoid-Pixel7_Neur200_Batch512_Dropout_Epoch250

By changing the activation function to sigmoid we see a clear improvement with a slightly lower score. The confusion matrix also improved, but those numbers are still too high. We can actually see a value of 30 which is not good.

Model:

- Activation function: tanh
- Neurons: 200
- Batch size: 1024
- Dropout: 0.5
- Epochs: 250
- Pixels: 7
- Orientations: 8

Test score: 0.1260688304901123
Test accuracy: 0.9575999975204468




Figure 21: ARN-HOG-Plot-tanh-Pixel7_Neur200_Batch1024_Dropout_Epoch250




Figure 22: ARN-HOG-ConfMat-tanh-Pixel7_Neur200_Batch1024_Dropout_Epoch250

We tried to increase the batch size which improved the curve for tanh activation function and lowered the gap between the two curves. The test set looks still too easy for the model. The confusion matrix shows really bad results. Numbers from 2 to 9 classification generates many errors.

Model:

- Activation function: sigmoid
- Neurons: 200
- Batch size: 1024
- Dropout: 0.5
- Epochs: 250
- Pixels: 7
- Orientations: 8




Figure 23: ARN-HOG-Plot-sigmoid-Pixel7_Neur200_Batch1024_Dropout_Epoch250




Figure 24: ARN-HOG-ConfMat-sigmoid-Pixel7_Neur200_Batch1024_Dropout_Epoch250

With the sigmoid function the results are the same.

Model:

- Activation function: tanh
- Neurons: 200
- Batch size: 1024
- Dropout: 0.5
- Epochs: 200
- Pixels: 7
- Orientations: 16

Test score: 0.10891010612249374
Test accuracy: 0.9648000001907349




Figure 25: ARN-HOG-Plot-tanh-Pixel7_Or16_Neur200_Batch1024_Dropout_Epoch200




Figure 26: ARN-HOG-ConfMat-tanh-Pixel7_Or16_Neur200_Batch1024_Dropout_Epoch200

Changing the number of orientations to 16 changed the curves with tanh function. It looks like a really good model because the curves are overlapping the entire time. Again, the confusion matrix shows something we couldn't see with the graph. This model is not really good at predicting what number it sees.

Model:

- Activation function: relu
- Neurons: 150
- Batch size: 1024
- Dropout: 0.5
- Epochs: 200
- Pixels: 7
- Orientations: 8

Test score: 0.09441999346017838
Test accuracy: 0.9700999855995178




Figure 27: ARN-HOG-Plot-relu-Pixel7_Or8_Neur150_Batch1024_Dropout_Epoch200




Figure 28: ARN-HOG-ConfMat-relu-Pixel7_Or8_Neur150_Batch1024_Dropout_Epoch200

Conclusion - HOG

We performed various tests and modified every parameter, we didn't find the perfect model. As we can see in most confusion matrices, the models tend to have difficulties with classifying numbers between 2 and 9. We see some values that stick out, such as 10, 22, 30, etc. The model {sigmoid, 200 neurons, 512 batch size, 0.5 dropout, 4 pixels, 9 orientations, 100 epochs} is quite good and it has the lowest amount of wrong classifications. The two curves meet at around 50 epochs which is pretty early compared to the other ones. In every graph we can see that the test curve is always below the training one. This shows that the test set might be too easy. The f1_scores we computed were good and around 0.95. As stated in the RAW part, we think the f1_score is not sufficient to determine if a model is good.

CNN.ipynb

Model:

- L4 neurons: 25
- L4 activation function: Relu
- Batch size: 2048
- Epochs: 50

Test score: 0.044601865112781525
Test accuracy: 0.9894000291824341




Figure 29: ARN-CNN-Plot-relu-Batch256_25L4_Epoch50




Figure 30: ARN-CNN-ConfMat-relu-Batch256_25L4_Epoch50

As we can see, the results with this preconfigured model are not that bad. The score is low and the accuracy is close to 1.0. However, from 10 epochs this model slowly starts to overfit. We can clearly see the gap between the training error and the testing error increases as we continue to iterate through the epochs.

Model:

- L4 neurons: 10
- L4 activation function: Relu
- Batch size: 2048
- Epochs: 50




Figure 31: ARN-CNN-Plot-relu-Batch256_10L4_Epoch50




Figure 32: ARN-CNN-ConfMat-relu-Batch256_10L4_Epoch50

We tried to decrease the number of neurons in the L4 layer. The scores are quiet similar to the previous model, but the testing error fluctuates more especially from the start to around 25 epochs. There's still

a small overfitting but the curves seem to be stable.

Model:

- L4 neurons: 5
- L4 activation function: Relu
- Batch size: 2048
- Epochs: 50




Figure 33: ARN-CNN-Plot-relu-Batch256_5L4_Epoch50






Figure 34: ARN-CNN-ConfMat-relu-Batch256_5L4_Epoch50

Here we tried an extremely low number of neurons in the L4 layer. The result is pretty good. The score is a little bit higher than the first model tested, but the two curves are almost overlapping the entire time. It looks like a good model. The confusion matrix shows that the number 4, 5 and 9 are quiet often wrongly classified.

Model:

- L4 neurons: 35
- L4 activation function: Relu
- Batch size: 2048
- Epochs: 50

**Figure 35:** ARN-CNN-Plot-relu-Batch256_35L4_Epoch50**Figure 36:** ARN-CNN-ConfMat-relu-Batch256_35L4_Epoch50

We then chose a high number of neurons to see how the model reacts. It's definitely not a good model

and 35 neurons is probably a bit too much for this task. There's an overfitting starting at 20 epochs. Surprisingly (or not), despite our bad curves the accuracy is the best we've had. The confusion matrix doesn't show a lot, the classification is not so bad.

Model:

- L4 neurons: 25
- L4 activation function: tanh
- Batch size: 2048
- Epochs: 50




Figure 37: ARN-CNN-Plot-tanh-Batch256_25L4_Epoch50

We also wanted to see how the preconfigured model would behave when changing the activation function in layer 4. This one uses tanh and we can clearly see it overfits directly at 10 epochs, maybe it's just bad luck and after tweaking some parameters it will probably show something good, but we didn't try since we were told just to change the number of neurons in the feed-forward part.

Model:

- L4 neurons: 25
- L4 activation function: sigmoid
- Batch size: 2048
- Epochs: 50




Figure 38: ARN-CNN-Plot-sigmoid-Batch256_25L4_Epoch50

With sigmoid the score is really low, the accuracy close to 1.0 and the curves are overlapping. After 20 epochs the testing error looks like it is going away a bit, but the gap is still really small after 50 epochs. The confusion matrix doesn't show abnormalities, there are few errors here and there, but nothing out of the ordinary.

Conclusion - CNN

As we can see lowering the number of neurons in L4 helped us reduce the gap between the testing and training curves with relu activation function. However, when reaching really low numbers like 5, the really low numbers of neurons, the confusion matrix shows something pretty different than the plot. Some numbers are not classified properly such as 4,5 and 9. Sometimes it represents 1-2% of the predictions for one particular numbers. It's not a big deal, but if we need to be really precise, this result is not enough. Other than that the digits are globally correctly classified even when changing the activation function. The f1_score we computed were around 0.95 with sometimes values up to 0.99.

Conclusion - CNN Fashion MNIST

For this experiment, we used the code provided in the CNN notebook and used the Fashion MNIST dataset provided by Keras. We then ran the model multiple times by changing some settings and watching the results. Below are some results that lead us to our final model.

Model:

```
1 batch_size = 16382
```

```

2 n_epoch = 100
3 Activation: relu
4
5 -----
6 Layer (type)                Output Shape                Param #
7 -----
8 l0 (InputLayer)              [(None, 28, 28, 1)]         0
9 l1 (Conv2D)                   (None, 28, 28, 9)           234
10 l1_mp (MaxPooling2D)         (None, 14, 14, 9)           0
11 l2 (Conv2D)                   (None, 14, 14, 9)           2034
12 l2_mp (MaxPooling2D)         (None, 7, 7, 9)             0
13 l3 (Conv2D)                   (None, 7, 7, 16)            1312
14 l3_mp (MaxPooling2D)         (None, 3, 3, 16)            0
15 flat (Flatten)               (None, 144)                  0
16 l4 (Dense)                    (None, 25)                   3625
17 l5 (Dense)                    (None, 10)                   260
18 -----
19 Total params: 7,465
20 Trainable params: 7,465
21 Non-trainable params: 0

```

Test score: 0.5433769822120667
Test accuracy: 0.7975999712944031




Figure 39: Model 1




Figure 40: Model 1

For this model, we used the same parameters as the CNN notebook. This model took a long time to be trained and the performances are bad with a test score of 0.54 and a Test accuracy of 0.79

Model:

Then we changed the activation function to tanh. The results were slightly better.

```
1 batch_size = 16382
2 n_epoch = 100
3 Activation: tanh
```

Layer (type)	Output Shape	Param #
l0 (InputLayer)	[(None, 28, 28, 1)]	0
l1 (Conv2D)	(None, 28, 28, 9)	234
l1_mp (MaxPooling2D)	(None, 14, 14, 9)	0
l2 (Conv2D)	(None, 14, 14, 9)	2034
l2_mp (MaxPooling2D)	(None, 7, 7, 9)	0
l3 (Conv2D)	(None, 7, 7, 16)	1312
l3_mp (MaxPooling2D)	(None, 3, 3, 16)	0
flat (Flatten)	(None, 144)	0
l4 (Dense)	(None, 25)	3625
l5 (Dense)	(None, 10)	260
Total params: 7,465		
Trainable params: 7,465		
Non-trainable params: 0		

Test score: 0.4326428771018982

Test accuracy: 0.8422999978065491




Figure 41: Model 2




Figure 42: Model 2

Model:

After a few try, and since it was taking a long time to train de models, we tried a different approach: reducing the model's complexity and removing features and adding neurons. We removed 2 layers of features extractions and adder neurons.

The model was a lot faster to train and the results were better with a test score of 0.32 and a test accuracy of 0.88.

```
1 batch_size = 8192
```

```

2 n_epoch = 50
3 Activation Function: relu
4
5 -----
6 Layer (type)                Output Shape          Param #
7 -----
8 l0 (InputLayer)              [(None, 28, 28, 1)]   0
9 conv2d (Conv2D)              (None, 28, 28, 64)    640
10 l1_mp (MaxPooling2D)         (None, 14, 14, 64)    0
11 flat (Flatten)              (None, 12544)          0
12 l4 (Dense)                   (None, 100)           1254500
13 l5 (Dense)                   (None, 10)            1010
14 -----
15 Total params: 1,256,150
16 Trainable params: 1,256,150
17 Non-trainable params: 0

```

Test score: 0.2910390794277191
Test accuracy: 0.8931999802589417




Figure 43: Model 3




Figure 44: Model 3

Final Model:

After a few tries with various parameters we finally landed of the following model, which was faster to train and yielded the best results. Our test set indicate a slight overfitting but stays stable.

- Test Score: 0.29
- Test accuracy: 0.89
- F-Score: 0.9
- One Conv2D layer: Conv2D(9, (3, 3)...)


```

1 batch_size = 4096
2 n_epoch = 100
3 Activation Function: relu
4
5 -----
6 Layer (type)                Output Shape                Param #
7 -----
8 l0 (InputLayer)              [(None, 28, 28, 1)]         0
9 l1 (Conv2D)                   (None, 28, 28, 9)           90
10 l1_mp (MaxPooling2D)          (None, 14, 14, 9)           0
11 flat (Flatten)               (None, 1764)                0
12 l4 (Dense)                   (None, 100)                 176500
13 l5 (Dense)                   (None, 10)                  1010
14 =====
15 Total params: 177,600
16 Trainable params: 177,600
17 Non-trainable params: 0

```

Test score: 0.2672470808029175
Test accuracy: 0.9054999947547913





Figure 45: Final Model

**Figure 46:** Final Model